

Continuity That Executes, Not Certifies

What a real BCP looks like when tested under load, drawn from moving a 450-person operation to work-from-home over a single weekend

A Field Note

July 2026 | Prepared by YuktiNexus

Contact, pradeep.sharma@yuktinexus.com

Confidential | For Operations, Risk, and Continuity Leaders | Distribute on Request

Two Different Documents Called the Same Thing

Every operation of any size has a business continuity plan, and most of those plans have passed an audit, a certification review, or a regulator's inspection at some point. Passing that review proves the document is complete, internally consistent, and aligned to a recognised framework. It does not prove the plan works, because a certification review tests the document against a standard. It does not, and cannot, test the document against an actual crisis, with real people, real systems, and real timing, none of which cooperate with what was written.

This note is about the difference between those two things, a continuity plan built to certify and a continuity plan built to execute, told through a specific weekend in which one operation's plan was tested by an actual, unplanned event rather than a scheduled drill. The plan had been certified twice. It had never been executed under real conditions. Both things turned out to be true at once, and the gap between them is the subject of this note.

The Weekend the Plan Actually Got Used

Friday Evening, the Trigger

The decision to move the operation to work-from-home arrived late on a Friday, with instruction to be fully remote by Monday morning, a compressed timeline the certified plan had, on paper, already accounted for. The documented plan specified device provisioning lead times, VPN capacity thresholds, and a phased rollout sequence, all of which had been reviewed and signed off in the most recent audit cycle. None of it had ever been run end to end, under time pressure, with the full population, before that Friday night.

Saturday, Where the Document and Reality Diverged

The first gap appeared within hours. The plan's device provisioning process assumed a steady-state request rate, a handful of laptop requests processed through the normal IT service desk queue, because that was the only rate anyone had ever tested it against. Four hundred and fifty simultaneous requests, arriving at once, did not behave like a scaled-up version of the normal process, it broke the process entirely, since the ticketing system's own workflow assumed requests would arrive with enough spacing for a human to triage each one individually.

The second gap was capacity, not process. VPN licensing had been provisioned to a documented headcount figure that was, it turned out, eighteen months out of date, reflecting the operation's size at the time of the last certification review rather than its actual current size. The plan had been certified against a number that was no longer true, and nobody had flagged the drift because certification checks the document's internal logic, not whether its inputs still match current reality.

FROM THE TRENCHES, A senior operations leader who has carried C-suite accountability across multi-tower captives

The number that was true eighteen months ago

We hit the VPN license ceiling by Saturday afternoon, at roughly seventy percent of the workforce connected. The plan said we had headroom for the full population. We didn't, because the plan's headcount figure was from the last time someone had updated the continuity documentation, and the operation had grown by nearly a hundred people since then without anyone circling back to check whether the BCP's numbers still matched. Nobody had done anything wrong, exactly, the documentation process had been followed correctly. It just hadn't been designed to notice that the underlying facts it depended on had quietly gone stale.

Saturday Night Into Sunday, the Real Plan Emerges

What actually got the operation to a working remote state by Monday was not the certified plan, it was a small group of people who understood the systems well enough to improvise around what the plan had gotten wrong, negotiating emergency license capacity, manually triaging the device request queue outside the broken ticketing workflow, and making real-time judgment calls the documented plan had no provision for because its authors could not have anticipated this specific combination of failures. This is not a criticism of those people, their judgment is exactly why the operation was live by Monday. It is a description of what the plan was actually missing, a means of executing under exactly this kind of pressure, not just a description of what should happen in principle.

Monday, Live but Not as Planned

The operation was fully remote and operational by Monday morning, which by any external measure looked like a continuity success, the plan had worked. Internally, everyone who had lived through the weekend knew the actual sequence of events looked nothing like the document, and that the difference between success and a genuine failure had come down to a small number of people improvising competently under pressure, not to the plan itself functioning as designed.

What the Certified Plan Actually Got Wrong

None of the gaps that surfaced that weekend were exotic or unforeseeable in the abstract. Each one is a recognisable category of continuity planning failure, and each had already passed a formal certification review without being caught, which is the more important point than any individual gap.

What the plan assumed	What actually happened	The underlying failure
Steady-state request processing scales linearly	450 simultaneous requests broke the workflow entirely, not just slowed it	The plan was never tested at the volume a real invocation would produce
Headcount and licensing figures were current	Both were roughly eighteen months stale	Certification checks internal consistency, not whether inputs still reflect reality
The documented process covers the response	The actual response was improvised by specific individuals working around the plan	The plan described an ideal sequence, not an executable one under real failure conditions

Why Certification and Execution Diverge

A Certification Review Tests the Document

A certification or audit review is, by design, a review of the document, its completeness against a framework, its internal logic, its alignment to policy and regulatory expectation. This is genuinely valuable and entirely insufficient on its own, because a document can be complete, internally consistent, and fully aligned to a recognised standard while still describing a process that has never been run under real conditions and would fail the first time it was.

A Drill Tests the Document Too, Just With People Reading From It

Scheduled continuity drills are a partial improvement, since they at least exercise the plan with real people rather than leaving it purely theoretical. Most scheduled drills, however, are run with advance notice, during business hours, with the people best equipped to handle a genuine crisis deliberately available and often already primed on what the drill will test. This produces a valuable but incomplete signal, closer to a rehearsal of the plan's happy path than a genuine test of whether it survives contact with an unplanned, poorly timed, and genuinely disruptive event.

Practitioner note, this is not an argument against certification or scheduled drills, both remain necessary. It is an argument that neither one, on its own, tells an organisation whether its continuity plan will actually work, and treating certification as proof of resilience is the specific mistake this note is trying to correct.

What a Plan Built to Execute Actually Requires

Test at the Volume a Real Invocation Would Produce

The device provisioning failure in this note's central example was invisible to certification because certification never asked the process to handle four hundred and fifty simultaneous requests. Any process embedded in a continuity plan needs to be tested, at least once, at something close to the volume a genuine full-scale invocation would actually generate, not at the steady-state volume the process handles on an ordinary day.

Treat the Plan's Inputs as Perishable

Headcount, licensing capacity, vendor contact details, and system inventories embedded in a continuity plan are only accurate at the moment they are written down, and they decay steadily afterward as the operation changes. A plan reviewed annually against a framework, without a parallel, more frequent process for refreshing its underlying data, will accumulate exactly the kind of quiet drift that turned a documented VPN headroom figure into a Saturday-afternoon capacity ceiling.

Run at Least One Genuinely Unannounced Test

- Schedule the announced drills that certification and stakeholder confidence require, and treat them as necessary but insufficient
- Separately, run at least one exercise a year without advance notice to the people who would actually be executing the response, even if it is narrower in scope than a full invocation
- Deliberately include a compressed or inconvenient timeline in that unannounced test, a weekend, a holiday period, a moment when the most experienced people are not readily available, since that is closer to how real invocations actually happen

FROM THE TRENCHES, A transitions leader who has executed lift-and-shift and Build-Operate-Transfer mandates end to end

The drill that never found the real gap

We'd run the same continuity drill for three years running, always on a Tuesday afternoon, always announced two weeks ahead, and it always passed cleanly. It wasn't until an actual weekend event forced an unplanned invocation that we found out the drill had never once tested what happens when the two people who know the failover process best are both unreachable at the same time. That specific gap had been sitting in the plan for three years, invisible to every drill we'd run, because every drill we'd run happened to occur when both of those people were sitting at their desks.

The Cost of Finding Out During a Real Event

The operation in this note's central example was fully functional by Monday morning, and that outcome depended on a small number of specific people being available, capable, and willing to improvise under significant pressure on a weekend. That is a genuinely fortunate outcome, not a repeatable one, since the entire recovery rested on individual judgment the certified plan had no mechanism to guarantee would be present the next time. An organisation that discovers its continuity gaps during a real invocation is relying on the same fortunate combination of circumstances every time, until, eventually, it isn't there.

The alternative is not a larger or more detailed planning document. It is closing the specific gap between what gets certified and what gets executed, testing at real volume, treating the plan's data as perishable, and deliberately finding the gaps before an actual weekend finds them first.

Designing a Test Programme That Finds Gaps Before Reality Does

Most continuity test programmes run a single type of exercise repeatedly, typically the announced, business-hours drill described earlier in this note, because it is the easiest to schedule, the least disruptive to run, and the version stakeholders and auditors most readily recognise. A test programme built to actually find gaps needs more than one type of exercise, each designed to surface a different category of failure.

Exercise type	What it is good at finding	What it misses
Announced, business-hours drill	Whether the documented process is logically complete and understood by the team	Volume failures, stale data, and gaps that only appear when key people are unavailable
Unannounced, limited-scope test	Whether the team can execute without advance preparation or rehearsal	Full-scale volume effects, since scope is usually deliberately narrow
Full-volume load test	Whether a specific process, such as device provisioning, survives real invocation-scale demand	Human and coordination failures, since it typically tests a system or process in isolation
Post-incident review of a real invocation	The most accurate possible picture of what actually breaks under genuine pressure	Nothing, but it only happens after the fact, and only for organisations that have already had a real event

The practical implication is that no single exercise type is sufficient on its own, and a mature test programme rotates deliberately through several of these categories across a planning cycle, rather than repeating the same announced drill because it is the version that has always been run. The specific mix matters less than the discipline of including at least one exercise per cycle designed specifically to surface the failure modes an announced drill structurally cannot.

Training for the Invocation, Not Just the Drill

A team that has only ever executed a continuity plan during scheduled, low-pressure drills has practised a different skill than the one a genuine invocation requires. The drill trains familiarity with the document. A real invocation demands improvisation around what the document gets wrong, coordination under time pressure, and decision-making with incomplete information, none of which a clean, well-rehearsed drill specifically builds. This is not an argument against drills, which remain useful for the skill they do build. It is an argument that an organisation relying on drill performance as a proxy for real-event readiness is measuring the wrong thing.

The teams that perform best during a genuine invocation are, consistently, the ones that have been deliberately exposed, at some point, to a exercise that did not go according to script, where something was deliberately withheld, delayed, or complicated beyond what the documented plan anticipated. Building at least one such exercise into the annual test programme is a direct, practical way to develop the specific capability a real invocation will actually demand.

A Quick Self-Audit

A small number of direct questions, asked honestly, usually reveal whether a continuity plan has actually been tested under conditions resembling a real invocation, or only reviewed against a certification framework.

- Has any process embedded in the plan ever been run at the full volume a genuine invocation would produce, or only at the scale of a routine, announced drill?

- When was the headcount, licensing, and system inventory data underlying the plan last verified against current reality, as distinct from when the document itself was last formally reviewed?
- Has the plan ever been invoked, even partially, without advance notice to the people expected to execute it?
- If the two or three people who know the failover and escalation process best were all simultaneously unreachable, does the plan specify what happens next, or does it implicitly assume they will be available?
- Has anyone traced, in writing, what actually happened during the last real invocation or major drill, against what the document said should happen, and closed the specific gaps that surfaced?

Practitioner note, a plan that cannot answer most of these questions with a confident yes has very likely been certified without ever being genuinely execution-tested, regardless of how recently or how well it passed its last formal review.

The Regulatory Reporting Angle

For regulated operations, particularly BFSI, business continuity is not only an internal resilience question, it is a documented, and in many cases regularly submitted, regulatory obligation, business impact analyses and continuity attestations filed with the relevant regulator on a defined cycle. The gap this note describes carries a specific additional risk in this context, an organisation can be in full, current, technically accurate regulatory compliance, having submitted every required BIA and continuity attestation on schedule, while still carrying exactly the kind of execution gap described in this note, invisible to the regulatory submission because that submission tests the same thing certification tests, the completeness and currency of the document, not its behaviour under real, unplanned load.

This is not a criticism of regulatory BIA processes, which serve a genuine and necessary purpose. It is a caution against treating regulatory compliance as a proxy for operational resilience, when the two are related but distinct, and an organisation that conflates them is exposed to precisely the failure mode this note opened with, a plan that is certified, compliant, and untested all at the same time.

About YuktiNexus

YuktiNexus is India's first practitioner expert network built specifically for operations and GCC transformation, connecting enterprises directly to practitioners who have built and run the mandate in question, not consultants who have studied it from outside. This field note is the fourth release in the YuktiNexus Practitioner Library, following Setting Up a GCC, The First 36 Months, When the Batch Window Is the Problem, and Pricing the Invisible, The White-Gloves Model.

Contact, pradeep.sharma@yuktinexus.com | +91 73307 46660 | www.yuktinexus.com