

Setting Up a GCC

The First 36 Months

The decisions made in the first three years that quietly determine the next ten

A White Paper

July 2026 | Prepared by YuktiNexus

Contact, pradeep.sharma@yuktinexus.com

Confidential | For GCC Sponsors, Operating Leaders, and Boards | Distribute on Request

Contents

Abstract

Executive Summary

Part 1, Orientation

Part 2, Months 1 to 6, Decide and Design

Section 1 Why GCCs Fail Before They Open

Section 2 Entity Structure and Site Selection

Section 3 Governance Design, Who Owns What Decision

Section 4 The Operating Model Decision, Captive, BOT, or Hybrid

Part 3, Months 6 to 18, Stand Up and Stabilise

Section 5 Leadership Hiring, The Hardest Decision to Get Right

Section 6 Talent Acquisition at Scale

Section 7 SLA Design, The Contract You Anchor On Day One

Section 8 Technology, Infrastructure, and the First Security Baseline

Section 9 The Compliance Foundation

Part 4, Months 18 to 36, Scale and Prove Value

Section 10 From Cost Centre to Value Centre

Section 11 Building the Second and Third Towers

Section 12 Culture at Scale, What Breaks Past 500 People

Section 13 The Retention Cliff

Section 14 Measuring What the Board Actually Wants

Part 5, Conclusions

Section 15 The 36-Month Roadmap, One Page

Section 16 What Comes After Month 36

Glossary

About YuktiNexus

Abstract

Most GCC setup guidance describes the org chart on day one. This paper describes the thirty-six months that follow it, the site decision nobody re-litigates until it is too late, the SLA anchored under time pressure that then governs every renewal for years afterward, the leadership hire that quietly decides whether the captive becomes a strategic asset or stays a cost line, and the attrition curve that arrives in month eighteen at exactly the moment the enterprise sponsor starts asking for savings.

This is not a market-sizing exercise. It draws on operations built and run from the ground up, both captive and outsourced, across BFSI, technology, and shared services, to document what the first three years actually require, and what happens when a decision made in month two goes unexamined until month twenty-four. The audience is anyone about to sponsor, lead, or govern a GCC for the first time, and anyone six months into one who suspects a foundational decision was made too quickly.

Executive Summary

Ten Findings That Define This Paper

1. The first ninety days decide more than the following thirty-four months combined.

Entity structure, site selection, and governance model are the three decisions made earliest and reversed latest. Every GCC that has struggled to scale past year two can trace at least one structural constraint back to a choice made in the first quarter, made under time pressure, by people who had not yet built the operation they were designing for.

2. Most GCCs are set up to be cost centres, and stay cost centres, because nobody designs the value narrative on day one.

The pivot from cost centre to value centre is treated as a year-three problem. It is a day-one design decision. The GCCs that make the transition build the measurement, reporting, and stakeholder narrative for it from the first quarter, not after the first efficiency target has already been hit and exhausted.

3. The SLA signed in month one becomes the contract you defend in year three.

Service levels anchored during setup, under pressure to show the parent organisation a signed document, routinely fail to reflect the operation that actually gets built. Renegotiating an SLA from a position of underperformance is a materially harder conversation than designing it correctly before volume arrives.

4. Leadership hiring is the highest-leverage, most commonly rushed decision in the entire setup.

The head of the GCC is hired to a timeline set by the parent's fiscal calendar, not by the actual difficulty of finding someone who can build a team, hold a board relationship, and run an operation simultaneously. Getting this hire wrong costs eighteen months to correct and is rarely fully correctable.

5. Attrition does not spike randomly. It spikes on a predictable calendar.

The eighteen-to-twenty-four month mark is when the first cohort has built enough market-visible skill to be actively recruited away, and it is also, almost without exception, when the parent organisation starts asking operations to demonstrate cost savings. Those two pressures arriving together is not a coincidence, it is structural, and it is predictable enough to plan for.

6. Governance ambiguity, not talent scarcity, is the most common cause of underperformance.

India has the talent. What GCCs run out of first is clarity about who owns which decision, between the parent function, the GCC leadership, and the steering committee. Most of the escalations that consume a GCC head's first year are governance failures wearing a talent-shortage costume.

7. The build, BOT, or hybrid decision is usually made on cost. It should be made on control and speed to value.

Cost comparisons between a captive build and a managed-services engagement routinely miss the value of speed, and the cost of losing direct control over the talent pipeline and institutional knowledge. The right model depends on what the enterprise actually needs to own, not on which option produces the lower number in year one.

8. Culture does not scale linearly. It breaks at specific headcount thresholds.

Roughly 150, roughly 500, and roughly 1,500 people each require a structurally different organisation, not simply more managers layered onto the same structure. GCCs that miss this treat every growth stage as a headcount problem instead of a design problem, and the culture erodes exactly when the operation needs it most.

9. Compliance debt accumulated in year one becomes expensive in year three.

The first ISO or SOC 2 surveillance audit typically lands in year two or three, and it tests every shortcut taken during the rush to go live. Retrofitting compliance into a live, revenue-critical operation costs several times what building it in from month one would have cost.

10. The GCCs that last ten years started designing for the cost-to-value pivot in month one, not year three.

By the time a GCC feels pressure to prove strategic value, the org design, the talent mix, and the stakeholder relationships already reflect three years of being run as a cost centre. Reversing that positioning is possible. It is dramatically harder than building toward it from the outset.

PART 1

ORIENTATION

The frame before the operational detail begins

Reading This Paper

This paper is organised chronologically, months one through thirty-six, because that is how the decisions actually arrive. It is written for four overlapping audiences, and each can read selectively without losing the thread.

The enterprise sponsor or board reader

Read the Executive Summary, Section 1, Section 3, and Section 15. This gives the strategic shape of the first three years without the operational detail underneath it.

The GCC head or COO, first mandate

Read the paper in full. Every section is written at the level of detail a first-time GCC leader needs, since this is the role most exposed to every decision documented here.

The HR, talent, and culture reader

Read Sections 5, 6, 12, and 13. These cover leadership hiring, talent acquisition at scale, culture thresholds, and the attrition curve in the detail this function needs.

The compliance, risk, and governance reader

Read Sections 3, 7, and 9. Governance design, SLA construction, and the compliance foundation are the three sections most directly relevant to this mandate.

Practitioner note, this paper draws on operational experience across BFSI, technology, and shared-services captives in India. Sector-specific regulatory requirements, particularly in BFSI and healthcare, add constraints beyond what is documented here, and should be validated against current regulatory guidance for the specific sector and geography involved.

PART 2

MONTHS 1 TO 6

Decide and design, the choices that are expensive to reverse

SECTION 1

Why GCCs Fail Before They Open

1.1 The Setup Timeline Is Set by the Wrong Clock

Nearly every GCC setup timeline is anchored to the parent organisation's fiscal calendar, a board commitment to be operational by a specific quarter, a budget cycle that needs the cost line to appear on schedule, or an executive sponsor who wants a visible win within their tenure. None of these clocks have any relationship to how long it actually takes to select a site, incorporate an entity, secure a facility, and hire a founding team capable of building an operation rather than simply staffing one.

The result is a familiar pattern. The timeline is fixed first, and every subsequent decision, site, entity structure, leadership hire, is compressed to fit it. Compression does not eliminate the work, it just moves the cost of doing it properly from month three, when it is cheap, to month eighteen, when it is expensive and now embedded in a live operation.

1.2 The Three Decisions That Get Made Too Fast

Three decisions consistently receive less diligence than they need, because they sit earliest in the timeline, before the operational pressure that would normally force rigor has arrived. Entity structure and site selection, covered in Section 2, governance design, covered in Section 3, and the build-versus-BOT decision, covered in Section 4. Each is treated as an administrative or procurement question. Each is, in practice, an operating-model decision with a ten-year horizon.

FROM THE TRENCHES, A C-level operator who built a GCC from the ground up

The site that looked identical on paper

We compared four cities using the standard scorecard, talent pool, cost, infrastructure, connectivity. Two came out nearly tied. We picked the one with marginally better real estate pricing. What the scorecard didn't capture was the depth of the specific skill mix we'd need by year two, not year one. By month eighteen we were running a parallel hiring pipeline in a second city because the first one had shallower bench strength in exactly the specialised roles that mattered most once we scaled past the generalist headcount. The scorecard wasn't wrong. It was measuring the wrong year.

1.3 The Founding Team Problem

A GCC's first fifty hires are disproportionately important, not because of their individual seniority, but because they become the institutional memory the next five hundred hires inherit. If the founding fifty are hired purely for immediate operational capacity, process execution, ticket resolution, transaction processing, the operation inherits a workforce culture built around execution rather than ownership. Retrofitting an ownership culture onto an execution-first workforce eighteen months later is materially harder than building it in from the first cohort.

Practitioner note, this does not mean over-hiring seniority in the early months. It means being deliberate about including, in the founding fifty, a meaningful proportion of people capable of designing process, not only running it, even if the immediate workload looks like pure execution.

Entity Structure and Site Selection

2.1 Entity Structure, the Choice That Outlives the People Who Made It

Whether the GCC is set up as a wholly owned subsidiary, a branch office, or through an intermediate holding structure determines tax treatment, repatriation mechanics, employment law exposure, and the ease of a future divestiture or restructuring, for the life of the entity. This decision is typically made by legal and tax advisors optimising for the immediate setup, with limited operational input, and it is rarely revisited once made, because unwinding an entity structure is materially more expensive than getting it right initially.

The practical failure mode is not choosing the wrong structure outright, most standard structures are defensible. It is failing to model how the structure behaves under scenarios the setup team is not thinking about in month two, a future acquisition of the GCC by a third party, a shift from single-parent captive to a shared multi-client model, or a significant headcount contraction. Structure decisions made purely for immediate tax efficiency routinely make these later scenarios harder than they needed to be.

2.2 Site Selection, Beyond the Standard Scorecard

Every site selection exercise compares broadly the same variables, talent availability, cost arbitrage, infrastructure quality, connectivity to the parent's time zones, and real estate cost. These variables matter, and most GCC setup teams evaluate them competently. What is more commonly under-evaluated is talent depth in the specific, specialised skills the operation will need at scale, not at launch.

A city can have an excellent general talent pool and a shallow bench in the three or four specialised capabilities that actually differentiate the GCC's value three years in, actuarial modelling, a specific regulatory compliance discipline, a niche engineering stack. Discovering this gap in year two means either building a second-city presence under time pressure, or accepting a permanent talent constraint on the operation's most strategic capabilities.

Evaluation dimension	What most scorecards measure	What is commonly missed
Talent pool	Aggregate graduate and lateral hire volume	Depth in the two or three specialised skills the operation will need at year three, not year one
Cost	Current salary benchmarks and real estate pricing	Cost trajectory as the same talent pool gets competed over by every other GCC entering the same city
Infrastructure	Connectivity, power reliability, transport access	Facility scalability, whether the same building or campus can absorb three years of headcount growth
Regulatory environment	State-level incentives and ease of incorporation	Sector-specific regulatory maturity for BFSI, healthcare, or other regulated verticals

2.3 The Multi-City Question

A growing number of GCCs now launch with a two-city model from month one, rather than adding a second city reactively once the first city's talent constraints surface. This costs more in setup complexity and duplicated governance overhead in year one. It buys optionality, talent diversification, and a business continuity posture that a single-city operation cannot match, at a point in the operation's life when adding that optionality is far cheaper than retrofitting it after a talent or continuity problem has already occurred.

Governance Design, Who Owns What Decision

3.1 The Governance Gap Nobody Notices Until It Costs Something

Ask any GCC head six months in what their single largest source of friction is, and the answer is rarely talent, infrastructure, or budget. It is ambiguity about who owns a given decision, the GCC's local leadership, the parent function's global process owner, or a steering committee that meets quarterly and is not positioned to make operational calls in real time.

This ambiguity is not a failure of any individual. It is the predictable result of designing an org chart before designing a decision-rights framework. Most setup projects produce a detailed reporting structure and no equivalent document specifying, for the ten or fifteen decision categories that recur monthly, who has final authority, who must be consulted, and who is simply informed.

3.2 A Minimum Viable Governance Framework

A functional governance framework does not need to be exhaustive on day one. It needs to cover the decisions that recur most often and cause the most friction when ownership is unclear, hiring approval above a certain level, process changes that affect SLA commitments, technology and tooling procurement, and escalation authority when a service level is at risk of being missed.

- Decision category, who has final authority, who must be consulted, who is informed, and the timeline within which the decision must be made
- A named escalation path with response-time commitments at each tier, not just an organisation chart
- A quarterly governance review that revisits the decision-rights framework itself, since the right ownership model at fifty people is rarely the right one at five hundred

3.3 The Steering Committee Trap

Steering committees are necessary and commonly misused. The failure mode is routing operational decisions, ones that need a same-week answer, into a body that meets monthly or quarterly. When this happens, the GCC's operating leadership either makes the call unilaterally and absorbs the political risk, or waits for the committee and absorbs the operational cost of the delay. Neither outcome is good, and both are avoidable with a governance framework that reserves the steering committee for strategic and escalated decisions specifically, not as the default path for anything ambiguous.

FROM THE TRENCHES, A senior operations leader who has carried C-suite accountability across multi-tower captives

The decision that took six weeks because nobody owned it

A process change that should have taken a week of approval sat for six because it touched three towers and none of the three tower heads believed it was their call to approve. It wasn't malice or incompetence, the decision genuinely sat in a gap between three reasonably-drawn organisational boundaries. We didn't fix it by redrawing the boundaries. We fixed it by adding a single rule, any decision that sits in a gap between two owners defaults to whichever owner is closer to the client relationship, with a right of appeal, not a requirement for consensus first.

The Operating Model Decision, Captive, BOT, or Hybrid

4.1 The Decision Framed as Cost, When It Should Be Framed as Control

Most build-versus-BOT comparisons are run as a cost exercise, total cost of a captive build against the fee structure of a managed-services or Build-Operate-Transfer arrangement. Cost matters, but a cost-only comparison misses the two variables that most determine whether the decision was right three years later, how much direct control the enterprise needs over the talent pipeline and institutional knowledge being built, and how much speed to operational value the enterprise actually needs.

A BOT model gets an operation to steady state faster, because the vendor brings an existing talent pipeline, established processes, and infrastructure already in place. What it costs is direct control during the build phase, and a transition event, at the point of transfer, that is itself a significant operational risk if not planned for from the start of the engagement, not treated as a year-three afterthought.

4.2 The Hybrid Model, More Common Than the Framing Suggests

In practice, few GCCs are purely captive or purely BOT. Most combine a captive core, the functions closest to the enterprise's competitive differentiation, with managed-services support for functions that are important but not differentiating, facilities management, certain infrastructure functions, or specific transactional processes at scale. The decision that matters is not choosing one model exclusively, it is being deliberate about which functions sit in which category, and revisiting that allocation as the operation matures rather than treating the initial split as permanent.

Model	Speed to steady state	Control retained	Where it fits best
Pure captive	Slowest	Highest	Functions core to competitive differentiation, long time horizon
Pure BOT	Fastest	Lowest during build	Functions needed at scale quickly, transfer plan is designed from day one
Hybrid	Moderate	Selective by function	Most enterprises, differentiated core plus supported non-core functions

4.3 The Transfer Event, If BOT Is Chosen

When a BOT model is selected, the transfer date is frequently treated as a milestone to be reached rather than an operational event to be engineered. A transfer executed without a structured knowledge-transfer period, without retention incentives for the vendor-side staff being absorbed, and without the receiving organisation's governance framework already in place, routinely produces a six-to-twelve month capability dip immediately after the point the enterprise expected to be operating at full strength.

Practitioner note, plan the transfer as its own project, with its own timeline, budget, and risk register, starting no later than twelve months before the planned transfer date, not as a line item inside the broader BOT contract.

PART 3

MONTHS 6 TO 18

Stand up and stabilise, the operation goes live and gets tested

SECTION 5

Leadership Hiring, The Hardest Decision to Get Right

5.1 A Different Role Than the Job Description Suggests

The GCC head role is frequently written and hired as an operations leadership position, someone who can run a large team and hit service levels. That is necessary and not sufficient. In the first eighteen months specifically, the role requires someone who can simultaneously build a team from a standing start, translate the parent organisation's strategic intent into an operating model that does not yet exist, and hold a credible relationship with a board or steering committee that has limited visibility into day-to-day operational reality.

These are three different skill sets, and it is uncommon to find all three at a strong level in a single candidate. Most GCC head searches optimise for the operational leadership skill, since it is the easiest to assess and the most obviously necessary, and under-weight the build and board-management dimensions, which are harder to evaluate in an interview process but matter just as much in the first eighteen months specifically.

5.2 What Actually Differentiates a Strong First GCC Head

- Direct prior experience building an operation from a standing start, not just running an established one, since the skills required to design a process are different from the skills required to optimise an existing one
- A demonstrated ability to make the case for investment upward, to a board or sponsor, using the language of enterprise value rather than only operational metrics
- Comfort operating with incomplete information and ambiguous governance in the first six to twelve months, since the governance framework in Section 3 will not be fully mature when this person starts

5.3 The Second-Layer Leadership Team

The GCC head cannot build the operation alone, and the second-layer leadership hires, the function heads reporting directly into the top role, receive materially less scrutiny than the top hire itself, despite carrying comparable long-term influence over how the operation actually runs. A weak second-layer hire in a critical function is harder to detect early, because the GCC head's own performance can mask it for the first year, and more expensive to correct later, because that person has by then built a team in their own image.

FROM THE TRENCHES, A transitions leader who has executed lift-and-shift and Build-Operate-Transfer mandates end to end

The hire that looked perfect on paper

We hired a function head with an outstanding résumé, exactly the pedigree the board wanted to see. Eight months in, the team under them had the highest attrition in the building. On paper the person was doing everything right, hitting every metric we tracked. What we weren't tracking was how they were getting there, through pressure rather than through building a team that wanted to stay. By the time attrition data made the problem visible, we'd lost eleven people from a team of forty, and half of our best process documentation left with them. We now interview explicitly for how a leadership candidate builds a team, not only for what they've delivered, and we track early-tenure attrition by manager, not only by function, from month one.

SECTION 6

Talent Acquisition at Scale

6.1 The Volume Problem Nobody Plans For

Hiring the first fifty people is a recruiting problem. Hiring the next four hundred and fifty, on a timeline that usually spans twelve to eighteen months, is a different kind of problem entirely, one of pipeline, employer brand, and process throughput that most GCCs are not staffed or structured to handle when the scaling phase begins. The recruiting function that worked well for a founding team of fifty routinely breaks somewhere between two hundred and four hundred hires, and the break is rarely visible until time-to-fill has already quietly doubled.

6.2 Employer Brand, Built Too Late

A new GCC entering an established market, a city with dozens of existing captives and outsourced operations competing for the same talent pool, starts with no employer brand recognition at all. Building that recognition takes months of consistent presence, campus relationships, and word-of-mouth reputation among the specific professional communities the operation needs to recruit from. GCCs that treat employer branding as a year-two marketing initiative, rather than a month-one operational necessity, pay for the gap in extended time-to-fill and a heavier reliance on higher-cost lateral hiring and agency fees throughout the scaling phase.

6.3 Quality at Volume, the Trade-off That Compounds

Under hiring-volume pressure, screening rigor is the variable that erodes first, quietly, because it is the hardest to measure in the moment and the easiest to justify relaxing when a requisition has been open for eight weeks. The cost of this erosion does not show up at the point of hire. It shows up eighteen months later, in the quality gap between the founding fifty and the four-hundredth hire, a gap that then has to be closed through remediation, additional training investment, or, in the more difficult cases, performance-managed exits that were themselves avoidable with tighter screening at the point of entry.

***Practitioner note,** build the recruiting function's capacity for volume before volume arrives, not in response to it. A recruiting team appropriately sized for fifty hires a quarter cannot simply absorb a target of two hundred without either a proportional increase in resourcing or a corresponding, and usually invisible, decline in quality.*

SECTION 7

SLA Design, The Contract You Anchor On Day One

7.1 Why the First SLA Is Almost Always Wrong

Service level agreements signed during GCC setup are typically negotiated before the operation exists to generate the data that should inform them. The GCC leadership team, under pressure to show the parent organisation a signed commitment, agrees to service levels based on benchmarks from other operations, vendor claims, or simple optimism, rather than on data from the specific team, process, and technology stack that will actually deliver against the commitment. This is not a failure of negotiation skill. It is a structural feature of doing SLA design before the operation is live.

7.2 The Renegotiation Problem

Renegotiating an SLA from a position of consistent underperformance is a fundamentally harder conversation than negotiating it correctly before volume arrives. It signals weakness to the parent organisation at exactly the point the GCC leadership needs credibility to make the case for additional investment, and it is frequently delayed

for exactly that reason, leaving the operation running for one or two additional years against a target it was never realistically going to hit, absorbing the reputational and morale cost of visible underperformance the entire time.

7.3 A Better Sequence

- Set provisional service levels for the first two to three months, explicitly labelled as provisional, with a scheduled review built into the original agreement rather than treated as a renegotiation
- Use the provisional period to gather real operational data, volume patterns, complexity distribution, and the actual capability of the team and technology as built, not as originally scoped
- Formalise the SLA at the scheduled review point, with the parent organisation's expectation set from the start that this is when the real commitment begins, which removes the stigma from what would otherwise look like a renegotiation

FROM THE TRENCHES, An infrastructure and technology leader who has run large platform implementations inside regulated environments
The SLA that assumed a system that didn't exist yet

The service level we signed assumed a level of system automation that was still eight months from being built. Nobody signed that number dishonestly, it was the number in the original business case, and the business case became the SLA without anyone re-checking whether the underlying system was actually live yet. We spent the next year explaining a shortfall every month to a steering committee that had every right to ask why. The fix wasn't complicated in hindsight, tie the SLA formalisation to a specific, named technology milestone, not to a calendar date chosen before the technology roadmap was finalised.

SECTION 8

Technology, Infrastructure, and the First Security Baseline

8.1 The Build-Versus-Inherit Decision

A GCC's technology stack is either inherited from the parent organisation, built new, or some combination of both, and this decision is frequently made by the parent's central IT function with limited input from the operational leadership who will actually run against it daily. Inherited systems bring the advantage of familiarity and existing integration, and the disadvantage of frequently being designed for the parent's home-market operating context, not for a new, distributed operation with different connectivity, compliance, and scale characteristics.

8.2 Security and Access, Designed for Fifty People, Used by Five Hundred

Access control and security architecture designed during the founding phase, when the team is small, trusted, and highly visible to leadership, does not automatically scale to a five-hundred-person operation with multiple shifts, several process towers, and a meaningfully higher inherent risk surface. The most common failure pattern is not a dramatic security incident. It is a slow accumulation of access-control debt, individually reasonable exceptions granted under operational pressure, that collectively leaves the operation with an access model nobody could fully explain by month twenty-four.

Practitioner note, build the security and access architecture for the headcount and complexity the operation will have at month twenty-four, not for the headcount it has at month three, even though this means the initial framework will feel over-engineered for a founding team of fifty.

8.3 Infrastructure Resilience, the Cost of Deferring It

Business continuity infrastructure, redundant connectivity, backup power, disaster recovery capability, is expensive to build and easy to defer when the operation is small and the immediate pressure is to get live. The deferral is rational in the moment and expensive later, both because retrofitting resilience into a live, revenue-

critical operation is materially more disruptive than building it in from the start, and because the parent organisation's tolerance for a continuity gap tends to drop sharply exactly as the GCC becomes operationally important enough for an outage to matter.

The Compliance Foundation

9.1 Compliance Debt, Invisible Until the First Audit

Every shortcut taken to hit a go-live date, a documented process step skipped, an access review deferred, a policy written but not yet operationalised, becomes compliance debt. It is invisible during the first year, because the first external audit, whether ISO 27001, SOC 2, or a sector-specific regulatory review, typically does not land until year two or later. By the time it does, the debt has usually compounded, since new hires have been trained on the shortcut version of the process, not the compliant one, and the gap between documented policy and actual practice has widened.

9.2 Building the Foundation Before the Deadline Forces It

The GCCs that handle their first external audit smoothly are, without exception, the ones that built compliance-relevant documentation and process discipline into the operation from month one, treating it as a standing operational requirement rather than a project to be run in the six months before an audit is scheduled. The six-month pre-audit scramble is a recognisable and avoidable pattern, and it is expensive not only in direct remediation cost but in the operational disruption of retrofitting documentation and evidence onto processes that have already been running, differently, for a year or more.

9.3 Sector-Specific Compliance, Layered On Top

BFSI and other regulated-sector GCCs carry compliance obligations beyond the general ISO and SOC 2 frameworks most operations plan around, data localisation requirements, sector-specific regulatory reporting, and business continuity obligations that may include mandated testing frequencies and regulator notification requirements. These obligations are frequently under-scoped during setup because they sit with the parent organisation's compliance function, which may not be closely involved in the GCC's day-to-day design decisions.

Practitioner note, involve sector-specific compliance and regulatory stakeholders in the setup design phase directly, not only at the point of the first audit or regulatory filing, since retrofitting sector-specific requirements into an already-live operation is one of the more expensive corrections documented in this paper.

PART 4

MONTHS 18 TO 36

Scale and prove value, the operation matures under pressure

SECTION 10

From Cost Centre to Value Centre

10.1 The Pivot That Arrives Whether the GCC Is Ready or Not

Somewhere between month eighteen and month thirty, nearly every GCC faces the same question from its parent organisation, in one form or another, what is this operation delivering beyond the cost arbitrage it was originally built for. GCCs that have spent their first eighteen months purely executing against an efficiency mandate, with no parallel investment in capability that could plausibly be called strategic, arrive at this question with little to say.

The GCCs that answer this question well are, without exception, the ones that began building the answer in month one, not month eighteen, by deliberately including higher-value capability, analytics, process design authority, a genuine centre-of-excellence function, inside the operating model from the start, even while the majority of early headcount was still focused on execution.

10.2 What the Pivot Actually Requires

- A small, protected allocation of senior capability from month one dedicated to capability-building and process improvement, not fully consumed by transactional volume, even when the volume pressure argues for redeploying them
- A measurement framework that captures value beyond cost savings from the outset, quality improvement, cycle-time reduction, risk reduction, so the data exists when the strategic conversation arrives, rather than needing to be constructed retroactively
- A deliberate narrative, communicated consistently to the parent organisation from early in the relationship, that frames the GCC as a capability investment rather than purely a cost lever, so the eventual pivot conversation is a continuation of an existing narrative, not a sudden repositioning

10.3 The Trap of Early Success

Somewhat counter-intuitively, GCCs that deliver strong cost efficiency in the first year are at higher risk of getting permanently typecast as a cost centre, precisely because the early efficiency numbers are so easy to celebrate and report upward. The parent organisation anchors on the efficiency narrative because it is the most visible, measurable win available in year one, and that anchor is difficult to shift later even once the GCC has genuinely developed higher-value capability.

Building the Second and Third Towers

11.1 Tower Expansion Is Not Just More of the Same

Adding a second or third functional tower to an established GCC, moving from a single finance-and-accounting operation, for instance, into a broader shared-services model spanning HR operations, procurement, or technology support, is frequently planned as a scaled repeat of the original setup. It is not. The second tower inherits an existing governance structure, an existing culture, and an existing leadership team's bandwidth, all of which shape the new tower's success as much as anything specific to the new function itself.

11.2 The Shared Services Integration Question

The central design decision in multi-tower expansion is how much shared infrastructure, facilities, HR, technology platform, common leadership layer, gets built once and shared across towers, versus how much each tower duplicates independently. Full sharing captures efficiency and creates a single point of failure and a slower decision path for each individual tower. Full independence protects each tower's speed and forfeits most of the efficiency case for building multiple towers under one roof in the first place. Most mature multi-tower GCCs land on a deliberate middle position, shared infrastructure and a shared governance framework, with functional leadership and process design kept independent per tower.

11.3 Sequencing the Expansion

The GCCs that expand well sequence the second tower's launch to begin after the first tower has reached operational stability, not simultaneously with the first tower's own scaling phase. Launching a second tower while the first is still absorbing its own attrition curve and governance maturation splits senior leadership attention at precisely the moment both towers need it most, and is a recognisable cause of both towers underperforming their individual, and individually achievable, targets.

Culture at Scale, What Breaks Past 500 People

12.1 Culture Does Not Scale Linearly

A culture that works at 150 people, built substantially on direct relationships with founding leadership and a shared sense of building something new, does not survive unchanged at 500, and the version that stabilises at 500 will not survive unchanged at 1,500. Each threshold requires a structurally different approach to communication, recognition, and decision-making, not simply a proportionally larger version of what worked before. GCCs that treat every growth stage as a headcount problem, solved by adding more managers into the existing structure, consistently see culture erode exactly when the operation most needs it to hold.

12.2 The 150-Person Threshold

Below roughly 150 people, most employees have direct or near-direct visibility into senior leadership, and culture propagates substantially through observed behaviour and direct relationship. Past this point, the majority of the workforce experiences leadership at a distance, through their direct manager and through communication that has to be deliberately designed rather than assumed to happen organically. GCCs that do not consciously redesign their communication and recognition mechanisms at this threshold see a documented, if often unattributed, dip in engagement in the two quarters following the crossing.

12.3 The 500-Person Threshold

Past roughly 500 people, a GCC typically operates across multiple shifts, multiple process towers, and often multiple physical floors or buildings, and the organisation stops being one culture and starts being several sub-cultures loosely bound by a shared brand and shared leadership. The task past this threshold is not maintaining one uniform culture, which is no longer realistic, but ensuring the sub-cultures share enough common ground, on core values and non-negotiable standards specifically, that they still function as one organisation rather than drifting into functionally separate operations that happen to share a logo.

12.4 The 1,500-Person Threshold

At scale beyond roughly 1,500 people, the GCC typically requires a dedicated internal communications and culture function, a formal manager-development pipeline, since the organisation can no longer rely on informally identifying good managers, and a much more deliberate approach to preserving any sense of the original founding culture, which by this point exists mainly as institutional memory held by an increasingly small proportion of long-tenured staff.

FROM THE TRENCHES, A C-level operator who built a GCC from the ground up

The town hall that stopped working

For the first two years, a monthly all-hands town hall was our single most effective culture tool. People asked hard questions directly, leadership answered directly, and everyone felt heard. Past about six hundred people, the same format started failing, not because engagement dropped, but because the room got too large for it to feel like a conversation anymore. It became a broadcast, and people stopped asking real questions in it. We didn't diagnose this quickly. We kept running the same format for another two quarters, wondering why engagement scores were slipping, before we realised the format itself, not the content, was the problem, and moved to a smaller, rotating-group structure that let the direct-conversation feeling survive past the size where the original format could carry it.

The Retention Cliff

13.1 Why Month Eighteen to Twenty-Four Is Predictable

Attrition in a maturing GCC is not evenly distributed across time. It concentrates in a recognisable window, roughly eighteen to twenty-four months from launch, and the concentration is structural rather than incidental. By this point, the founding and early cohorts have built enough market-visible skill and experience to be actively and specifically recruited by competing operations in the same city, and the parent organisation has typically begun asking the GCC to demonstrate cost efficiency, which frequently translates into compensation and promotion constraints at exactly the moment retention pressure is highest.

13.2 The Compounding Cost

Attrition at this stage costs more than the direct replacement cost suggests, because the people leaving are disproportionately the ones who hold the institutional knowledge built during the founding phase, the informal process fixes, the undocumented client relationship context, the reasons certain decisions were made a certain way. Losing this knowledge does not show up as a line item. It shows up as a slow degradation in process quality and client relationship continuity that is difficult to attribute directly back to the attrition event that caused it.

13.3 What Actually Moves the Needle

- Building a genuine career path with visible, credible next roles before the eighteen-month mark arrives, not in response to the first wave of resignations
- Protecting compensation review cycles from being the first target of a parent-driven cost efficiency push, since this is precisely the wrong moment to compress the retention lever that matters most
- Deliberately over-documenting process and institutional knowledge before the eighteen-month mark, on the assumption that some proportion of the people holding it will leave regardless of retention efforts

Practitioner note, no retention programme fully eliminates this attrition window. The realistic goal is reducing its severity and ensuring the institutional knowledge at risk has been captured before the people holding it depart, not preventing the departures entirely.

Measuring What the Board Actually Wants

14.1 The Metrics Gap Between Operations and the Board

Most GCCs report operationally, service levels, ticket volumes, cost per transaction, because these are the metrics the operating team naturally tracks to run the business day to day. A board or steering committee, however, is typically evaluating a different question, whether the GCC investment is delivering value proportional to its cost and strategic positioning, and operational metrics alone rarely answer that question in the language a board is asking it.

14.2 A Layered Reporting Model

The GCCs that manage the board relationship well typically maintain two layers of reporting, an operational layer used internally to run the business, and a strategic layer, translated specifically for board and sponsor audiences, that connects operational performance to enterprise value, cost avoided, risk reduced, capability built, strategic optionality created. Building the translation capability between these two layers is a specific skill, and it is worth treating as a distinct function rather than assuming operational leaders will develop it naturally.

14.3 Leading Indicators, Not Just Lagging Ones

Most GCC reporting is heavily weighted toward lagging indicators, service levels achieved last month, cost delivered last quarter. Leading indicators, attrition risk signals, pipeline health for critical skills, the gap between documented process and actual practice, give the board and the operating leadership earlier warning of problems that would otherwise only become visible once they have already affected a lagging metric. Building a small set of genuinely predictive leading indicators, rather than simply more lagging ones, is one of the more differentiating investments a maturing GCC can make.

PART 5

CONCLUSIONS

Where the first thirty-six months lead, and what comes after

SECTION 15

The 36-Month Roadmap, One Page

The following summarises the paper's chronology as a single reference. It is not a substitute for the detail in the preceding sections, and it should be read alongside them, not instead of them.

Phase	Months	The decisions that matter most
Decide and design	1 to 6	Entity structure and site selection (Section 2), governance framework (Section 3), operating model, captive, BOT, or hybrid (Section 4)
Stand up and stabilise	6 to 18	Leadership hiring (Section 5), talent acquisition at scale (Section 6), SLA formalisation (Section 7), technology and compliance foundation (Sections 8 to 9)
Scale and prove value	18 to 36	The cost-to-value pivot (Section 10), tower expansion (Section 11), culture thresholds (Section 12), the retention cliff (Section 13), board-level reporting (Section 14)

SECTION 16

What Comes After Month 36

16.1 The Operation That Emerges

A GCC that has navigated the first thirty-six months deliberately, rather than reactively, arrives at month thirty-six with a materially different profile than one that has simply survived the same period. It has a governance framework that has been tested and revised at least once, a leadership bench two layers deep rather than dependent on the founding GCC head alone, a compliance posture built in rather than retrofitted, and, most importantly, a value narrative with the board that was constructed from month one rather than assembled defensively in month twenty-four.

16.2 The Decade Question

Every decision documented in this paper is, in some form, a bet on what the GCC needs to look like not at month thirty-six but at year ten. Entity structures chosen for immediate tax efficiency, sites chosen for year-one cost, and SLAs signed for political convenience are all, in effect, short-horizon decisions made inside what is structurally a long-horizon commitment. The GCCs that perform best over a decade are, consistently, the ones whose founding leadership treated the first thirty-six months as the design phase for the following ten years, not as a separate, self-contained project with its own success criteria.

This is the single throughline across every section of this paper. The first thirty-six months are not a runway to steady state. They are the design phase for everything the operation will become, and very little of what gets built during them is genuinely temporary, however temporary it may feel at the time.

Glossary

Key terms and abbreviations used throughout this paper.

Term	Definition
Attrition cliff	The predictable concentration of voluntary attrition in a GCC's eighteen-to-twenty-four month window, driven by skill maturity meeting parent-driven cost pressure simultaneously.
BOT (Build-Operate-Transfer)	An operating model in which a third-party vendor builds and initially operates a GCC on the enterprise's behalf, transferring it to direct enterprise control at an agreed future date.
Captive	A GCC wholly owned and operated by the parent enterprise, as distinct from an outsourced or BOT arrangement.
Centre of excellence (CoE)	A specialised capability within a GCC focused on process design, innovation, or a specific technical discipline, as distinct from transactional execution capacity.
Compliance debt	The accumulated gap between documented policy and actual operational practice, created by shortcuts taken to meet a launch or delivery timeline, which surfaces at the first external audit.
Cost centre	An organisational unit evaluated primarily on the cost efficiency it delivers, as distinct from a value centre evaluated on broader strategic contribution.
Founding cohort	The first substantial group of hires in a new GCC, typically the first fifty to one hundred people, who disproportionately shape the operation's early culture and institutional knowledge.
GCC (Global Capability Centre)	An offshore or nearshore operational unit established by an enterprise to deliver a defined set of business, technology, or knowledge functions on a captive or semi-captive basis.
Governance framework	A defined structure specifying decision rights, escalation paths, and review cadence between a GCC's operating leadership, the parent organisation, and any steering or governance body.
Hybrid model	An operating model combining a captive core with managed-services or BOT support for selected functions, rather than a purely captive or purely outsourced structure.
Institutional knowledge	Undocumented process context, client relationship history, and operational judgment held informally by tenured staff, at risk of loss during periods of high attrition.
Leading indicator	A metric that provides early warning of a future operational or strategic problem, as distinct from a lagging indicator that reports a result after it has occurred.
Second-layer leadership	The tier of functional leaders reporting directly to the GCC head, whose hiring quality has comparable long-term influence on the operation as the top hire itself.
SLA (Service Level Agreement)	A formal commitment specifying the performance standards a GCC will deliver against, typically covering timeliness, quality, and availability metrics.
Steering committee	A governance body, typically composed of parent-organisation and GCC leadership, responsible for strategic oversight and escalated decision-making.
Tower	A distinct functional service line within a GCC, such as finance and accounting, HR operations, or technology support, often added sequentially as the operation matures.
Value centre	An organisational unit evaluated on its broader strategic contribution to the enterprise, beyond direct cost efficiency, including capability, risk reduction, and innovation.

About YuktiNexus

YuktiNexus is India's first practitioner expert network built specifically for operations and GCC transformation. The network connects enterprises making a GCC, managed-services, or transformation decision directly to practitioners who have already built and run exactly that mandate, not consultants who have studied it from the outside.

The network was founded by Pradeep Sharma, who spent over two decades inside operations, building them from the ground up, the floor, the team, the culture, the systems, across 154 global markets and twenty-five outsourced client businesses, with direct experience in GCC setup and governance and in AI transformation inside regulated environments.

Five additional practitioners are networked and available alongside Pradeep, each carrying leadership-level accountability through the full life of a GCC, a C-level operator who built a GCC from entity incorporation and site selection through the first years of live, at-scale operations, an infrastructure and technology leader who has run large platform implementations and core system migrations inside regulated environments, an AI, ML, and blockchain practitioner who has taken these from proof of concept into production inside live financial operations, a transitions leader who has executed lift-and-shift and Build-Operate-Transfer mandates end to end, and a senior operations leader who has carried C-suite accountability across multi-tower captives.

Engaging YuktiNexus

Enterprises can engage the network across five formats, from a structured sixty-to-ninety minute Expert Call through a multi-day Validation Sprint, an embedded four-to-twelve week advisory project, fractional leadership placement for six to twelve months, or a structured Mentoring Circle for mid-career operations managers. Every match is made by hand and screened for conflict of interest before either party sees the other.

This paper is the first release in the YuktiNexus Practitioner Library, a series of white papers, playbooks, and field notes written by practitioners who built and ran these operations, not by researchers who studied them.

Contact, pradeep.sharma@yuktinexus.com | +91 73307 46660 | www.yuktinexus.com